

pytorch-vgg16

September 5, 2021

1 Synopsis

This project aims to build a deep learning model using PyTorch to classify images using convolutional neural networks. The model is trained on the well-known CIFAR-10 dataset (<https://www.cs.toronto.edu/~kriz/cifar.html>). The techniques of transfer learning (from the VGG16 model with batch normalization) and data augmentation were used to enhance the model's accuracy. A final accuracy of 87.8600 % was obtained on the test set.

2 Setup

Install the torchinfo library to obtain model summaries.

```
[ ]: ! pip install torchinfo
```

Import the libraries and methods needed for the project.

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import torch
from torch.nn import (
    Sequential,
    Linear,
    ReLU,
    Dropout,
    CrossEntropyLoss
)
from torchvision.transforms import (
    Compose,
    Resize,
    RandomRotation,
    RandomHorizontalFlip,
    RandomCrop,
    ToTensor,
    Normalize
)
from torchvision.datasets import CIFAR10
from torch.utils.data import (
```

```

        random_split,
        DataLoader
    )
    from torchvision.models import vgg16_bn
    from torchinfo import summary
    from torch.optim import Adam
    from copy import deepcopy

```

3 Preprocessing

Lay out the image size, means and standard deviations expected by the pretrained model.

```

[ ]: pretrained_size = 224
pretrained_means = [0.485, 0.456, 0.406]
pretrained_stds = [0.229, 0.224, 0.225]

```

Define transforms for the training set.

```

[ ]: train_transforms = Compose([
    Resize(size = pretrained_size),
    RandomRotation(degrees = 5),
    RandomHorizontalFlip(p = 0.5),
    RandomCrop(size = pretrained_size,
                padding = 10),
    ToTensor(),
    Normalize(mean = pretrained_means,
              std = pretrained_stds)
])

```

Define transforms for the test set.

```

[ ]: test_transforms = Compose([
    Resize(size = pretrained_size),
    ToTensor(),
    Normalize(mean = pretrained_means,
              std = pretrained_stds)
])

```

4 Load the data

Load the data, performing the predefined transforms.

```

[ ]: ROOT = ".data"

raw_train_data = CIFAR10(root = ROOT,
                          train = True,
                          download = True,

```

```

        transform = train_transforms)

test_data = CIFAR10(root = ROOT,
                    train = False,
                    download = True,
                    transform = test_transforms)

```

5 Data partitioning

Carve out a validation set from the raw training set. 90 % of the raw training set forms the final training set and the other 10 % forms the validation set.

```

[ ]: SPLIT_FRACTION = 0.9

n_train_examples = int(len(raw_train_data) * SPLIT_FRACTION)
n_valid_examples = len(raw_train_data) - n_train_examples

train_data, valid_data = random_split(dataset = raw_train_data,
                                     lengths = [n_train_examples,
                                     ↪n_valid_examples])

```

Ensure that the validation set uses the same transforms as the test set.

```

[ ]: valid_data = deepcopy(x = valid_data)
      valid_data.dataset.transform = test_transforms

```

Check the number of examples in each set.

```

[ ]: print(f"Number of examples in the training set: {len(train_data)}")
      print(f"Number of examples in the validation set: {len(valid_data)}")
      print(f"Number of examples in the test set: {len(test_data)}")

```

6 Check transforms

Plot 25 sample images to check whether the proposed transforms are sensible.

```

[ ]: def normalize_image(image):

      image_min = image.min()
      image_max = image.max()
      image.clamp_(min = image_min,
                  max = image_max)
      image.subtract_(image_min).div_(image_max - image_min + 1e-5)

      return image

```

```
[ ]: def plot_images(images,
                    labels,
                    classes,
                    normalize = True):

    num_images = len(images)

    rows = int(np.sqrt(num_images))
    columns = int(np.sqrt(num_images))

    fig = plt.figure(figsize = [10, 10])

    for i in range(rows * columns):

        ax = fig.add_subplot(rows, columns, (i+1))

        image = images[i]

        if normalize:
            image = normalize_image(image)

        ax.imshow(image.permute(1, 2, 0).cpu().numpy())
        ax.set_title(classes[labels[i]])
        ax.axis("off")
```

```
[ ]: N_IMAGES = 25
```

```
[ ]: images, labels = zip(*[(image, label) for (image, label) in
                           [train_data[i] for i in range(N_IMAGES)]])

classes = test_data.classes

plot_images(images, labels, classes)
```

7 Form data iterators

Form iterators for the training, validation and test sets using a desired batch size.

```
[ ]: BATCH_SIZE = 128
```

```
[ ]: train_iterator = DataLoader(dataset = train_data,
                               shuffle = True,
                               batch_size = BATCH_SIZE)

valid_iterator = DataLoader(dataset = valid_data,
                           batch_size = BATCH_SIZE)
```

```
test_iterator = DataLoader(dataset = test_data,
                           batch_size = BATCH_SIZE)
```

8 Define the model

Import a pretrained VGG16 model with batch normalization.

```
[ ]: model = vgg16_bn(pretrained = True)
```

Inspect the layers present in the model.

```
[ ]: model
```

Ensure that the convolutional feature-extracting base is frozen and the classifier head is unfrozen. Setting `requires_grad` to `False` freezes the corresponding layer and setting it to `True` unfreezes it. The adaptive average pooling layer contains no parameters.

```
[ ]: for param in model.features.parameters():
      param.requires_grad = False
```

```
[ ]: for param in model.classifier.parameters():
      param.requires_grad = True
```

The pretrained model was trained on the ImageNet dataset, which had 1000 classes. Hence, the final layer in the classifier has 1000 output features. Inspect this last layer.

```
[ ]: model.classifier[-1]
```

Since I am currently building a model on the CIFAR-10 dataset, which has 10 classes, I want this final layer to have 10 output features. Replace the final layer with a linear layer which has 10 output features. The newly added layer will have `requires_grad` set to `True` and will be trainable.

```
[ ]: N_CLASSES = 10
      IN_FEATURES = model.classifier[-1].in_features
```

```
[ ]: final_layer = Linear(in_features = IN_FEATURES,
                        out_features = N_CLASSES)
```

```
[ ]: model.classifier[-1] = final_layer
```

Check that the modified classifier does have 10 output features.

```
[ ]: model.classifier
```

Summarize the final overall model.

```
[ ]: summary(model,
            input_size = [BATCH_SIZE, 3, 224, 224],
            device = "cuda")
```

9 Loss function and optimizer

Define the loss function and the optimizer to be used.

```
[ ]: loss_fn = CrossEntropyLoss()
optimizer = Adam(params = model.parameters(),
                 lr = 0.001)
```

10 Copy the model to the GPU

```
[ ]: if torch.cuda.is_available():
    device = torch.device("cuda")
else:
    device = torch.device("cpu")
```

```
[ ]: print(f"Using {device} device")
```

```
[ ]: model.to(device)
```

11 Define the training function

Define a function to train the model and simultaneously validate it, across a desired number of epochs.

```
[ ]: def train(train_iterator,
              valid_iterator,
              model,
              loss_fn,
              optimizer,
              device,
              n_epochs = 5):

    # Initialize the history list which will contain all losses and metrics

    history_list = []

    for i in range(n_epochs):

        # Start setting up the training procedure

        print(f"Epoch {i+1}")
```

```

print("-----")

train_size = len(train_iterator.dataset)
n_train_batches = len(train_iterator)
train_loss = 0
average_train_loss = 0
train_n_correct = 0
train_accuracy = 0

# Set the model to training mode

model.train()

for train_batch, (X, y) in enumerate(train_iterator):

    # Copy the tensors to the GPU

    X = X.to(device)
    y = y.to(device)

    # Reset the gradients of the model parameters to zero

    optimizer.zero_grad()

    # Obtain the model prediction and loss

    pred = model(X)
    loss = loss_fn(pred, y)

    # Backpropagate the loss and deposit each gradient in place

    loss.backward()

    # Adjust the parameters using the gradients collected in the
    ↪ backward pass

    optimizer.step()

    # Increment the validation loss and the number of correctly labeled
    ↪ instances

    # Build up these aggregate values instance by instance

    train_loss += loss.item()
    train_n_correct += (pred.argmax(1) == y).type(torch.float).sum().
    ↪ item()

    # Display the training loss after every hundredth batch is trained

```

```

        if train_batch % 100 == 0:
            loss = loss.item()
            current_instance = train_batch * len(X)
            print(f"Loss: {loss:.6f} [{current_instance:5f} / {train_size:
→5f}]]")

    # Obtain average training loss and accuracy for the entire epoch

    average_train_loss = train_loss / n_train_batches
    train_accuracy = train_n_correct / train_size

    # After training is finished, start validation

    valid_size = len(valid_iterator.dataset)
    n_valid_batches = len(valid_iterator)
    valid_loss = 0
    average_valid_loss = 0
    valid_n_correct = 0
    valid_accuracy = 0

    with torch.no_grad():

        # Set the model to evaluation mode

        model.eval()

        for X, y in valid_iterator:

            # Copy the tensors to the GPU

            X = X.to(device)
            y = y.to(device)

            # Obtain the model prediction and loss

            pred = model(X)
            loss = loss_fn(pred, y)

            # Increment the validation loss and the number of correctly
→labeled instances
            # Build up aggregate values instance by instance

            valid_loss += loss.item()
            valid_n_correct += (pred.argmax(1) == y).type(torch.float).
→sum().item()

```



```

        # Obtain average validation loss and accuracy for the entire epoch

        average_valid_loss = valid_loss / n_valid_batches
        valid_accuracy = valid_n_correct / valid_size

        print("Validation error:")
        print(f"Accuracy: {valid_accuracy:.6f}, Average loss:␣
→{average_valid_loss:.6f}")
        print()

        history_list.append([average_train_loss, average_valid_loss,␣
→train_accuracy, valid_accuracy])

        # Display a message indicating training has finished

        print()
        print("Done!")

        # Create a data frame containing the entire training and validation history

        history = pd.DataFrame(data = history_list,
                                columns = ["average_train_loss",␣
→"average_valid_loss",
                                "train_accuracy", "valid_accuracy"])

        return model, history

```

12 Train the model

Pick a suitable number of epochs.

```
[ ]: n_epochs = 20
```

Run the training function.

```
[ ]: model, history = train(train_iterator = train_iterator,
                             valid_iterator = valid_iterator,
                             model = model,
                             loss_fn = loss_fn,
                             optimizer = optimizer,
                             device = device,
                             n_epochs = n_epochs)
```

Plot the training and validation loss.

```
[ ]: plt.figure(figsize = [10, 8])
      plt.plot(history["average_train_loss"],
```

```

        label = "Average training loss")
plt.plot(history["average_valid_loss"],
        label = "Average validation loss")
plt.legend()
plt.xlabel("Epoch")
plt.ylabel("Average loss")
plt.title("Loss plot")
plt.show()

```

Plot the training and validation accuracy.

```

[ ]: plt.figure(figsize = [10, 8])
plt.plot(history["train_accuracy"],
        label = "Training accuracy")
plt.plot(history["valid_accuracy"],
        label = "Validation accuracy")
plt.legend()
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.title("Accuracy plot")
plt.show()

```

13 Evaluate the model

Perform a final evaluation of the model on the test set. First, define a function to carry out the same.

```

[ ]: def evaluate(test_iterator,
                model,
                loss_fn,
                device):

    size = len(test_iterator.dataset)
    num_batches = len(test_iterator)
    test_loss = 0
    correct = 0

    with torch.no_grad():
        for X, y in test_iterator:
            X = X.to(device)
            y = y.to(device)
            pred = model(X)
            test_loss += loss_fn(pred, y).item()
            correct += (pred.argmax(1) == y).type(torch.float).sum().item()

    average_loss = test_loss / num_batches
    accuracy = correct / size

```

```
print("Test error:")  
print(f"Accuracy: {accuracy:.6f}, Average loss: {average_loss:.6f}")
```

Then, run the function on the model which has just been trained.

```
[ ]: evaluate(test_iterator = test_iterator,  
             model = model,  
             loss_fn = loss_fn,  
             device = device)
```

14 Save the model

Save the model to disk.

```
[ ]: torch.save(model, "pytorch-vgg16-cifar-10-model.pth")
```